



# Installing containerized software on Alpine

# Installing software on Alpine with Conda and Mamba

Date: February 27, 2024

Instructor: Andrew Monaghan

Contributors: Layla Freeborn, Trevor Hall, Brandon Reyes

- Website: [www.rc.colorado.edu](http://www.rc.colorado.edu)
- Documentation: <https://curc.readthedocs.io>
- Helpdesk: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Survey: <http://tinyurl.com/curc-survey18>

## Slides

[https://github.com/ResearchComputing/containerized\\_software\\_alpine\\_primer](https://github.com/ResearchComputing/containerized_software_alpine_primer)



# Learning Objectives

- What is a software container?
- How can I find existing software containers?
- How can I use a container on Alpine?
- How can I build a container on Alpine?

# Session Overview

## Introduction

- Installing software on CURC systems
- Defining containers
- Description of Apptainer

## Setting up Conda on Alpine

- Logging in
- Using conda for the first time: creating the ~/.condarc file
- Starting an interactive session and activating conda

## Creating and Modifying Virtual Environments with Conda

- Creating/activating/modifying a python environment
- Useful conda commands and paths

## Using Conda Virtual Environments

- In HPC jobs
- In OnDemand Jupyter

## Strategies for installing complex Virtual Environments (Discussion only)

- Channels
- Resolving conflicts upon environment creation
- Mamba

# Building Software on Alpine

- There are numerous ways to install software on Alpine:
  - grab pre-compiled binaries
  - compile from source
  - within virtual environments (via Conda, Miniconda, or Mamba)
  - **using containers (Apptainer)**
  - using a package manager for HPC systems (Spack)

*Additional information:*



<https://github.com/ResearchComputing/research-software-curc>



# What is a Container?

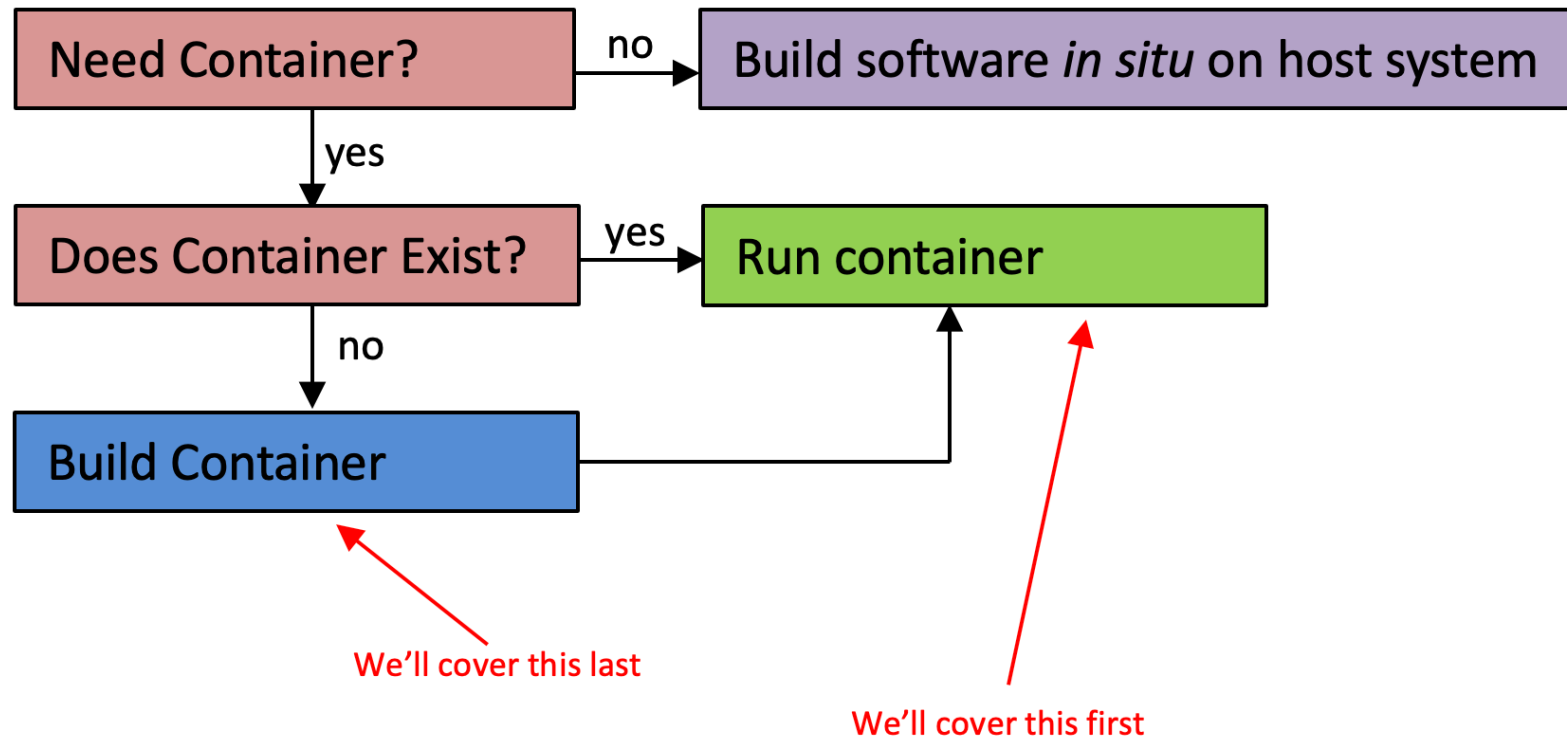
- A container is a portable environment that packages some or all of the following: an operating system, software, libraries, compilers, data, and workflows. Containers enable:
  - Mobility of compute
  - Reproducibility (software and data)
  - User Freedom

# Containerization Software

- Docker 
  - Well-established
  - Most widely used
  - Millions of containers already available on [DockerHub](https://hub.docker.com/)
- **Apptainer** 
  - Formerly “Singularity”
  - HPC-safe
- Others



# Making the decision to containerize



# Logging into CU Research Computing

login to CURC via your terminal:

```
ssh monaghaa@login.rc.colorado.edu
```

...or login to CURC via your browser:

<https://ondemand-rmacc.rc.colorado.edu>

(once logged in, navigate to **Clusters** -> **Alpine shell**)

*Additional information:*

<https://curc.readthedocs.io/en/latest/access/logging-in.html>

<https://curc.readthedocs.io/en/latest/gateways/OnDemand.html>

# Start a compute session

Start a session on an Alpine compute node with **acompile**:

```
[monaghaa@login11 ~]$ acompile --help
[monaghaa@login11 ~]$ acompile --ntasks=4 --time=90:00
...
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer --help
```

**Note:** when you login to CURC you'll be on a **login** node. You'll need to be on a **compute** node to use apptainer. The **acompile** command allows you to quickly start an interactive job on a compute node.

*Additional information:*

<https://curc.readthedocs.io/en/latest/clusters/alpine/alpine-hardware.html#partitions>

# Let's run our first container!

Invoke “R” statistical software in a pre-existing container:

```
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer exec $CURC_CONTAINER_DIR/debian_r_4_2_2.sif R  
> 1+2  
> quit ()
```

```
[monaghaa@c3cpu-a5-u28-1 ~]$ ls $CURC_CONTAINER_DIR/    # ← what the heck is in this directory?
```

Let's run an external R script with the container

```
[monaghaa@c3cpu-a5-u28-1 ~]$ cd /scratch/alpine/$USER  
[monaghaa@c3cpu-a5-u28-1 ~]$ git clone https://github.com/ResearchComputing/containerized\_software\_alpine\_primer  
[monaghaa@c3cpu-a5-u28-1 ~]$ cd containerized_software_alpine_primer/examples  
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer exec $CURC_CONTAINER_DIR/debian_r_4_2_2.sif Rscript prime_functions.R
```

# Key Apptainer Commands

[exec](#): Execute a command to your container

[run](#): Run your image as an executable (behavior must be predefined)

[build](#): Build a container on your user endpoint or build environment

[pull](#): pull an image from Docker or Singularity Hub

[inspect](#): See labels, run and test scripts, and environment variables

[shell](#): Shell into your image



# Now let's explore a container

Shell into the container:

```
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer shell $CURC_CONTAINER_DIR/gmtsar_v6.2.sif
Apptainer> cat /etc/os-release # what operating system is in the container?
Apptainer> ls /usr/local # what do you see?
Apptainer> exit
```

Inspect how the container was built:

```
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer inspect -d $CURC_CONTAINER_DIR/gmtsar_v6.2.sif
```

...now let's look more closely at a container definition file

# Container definition file

lolcow.def

```
Bootstrap: docker
From: ubuntu:20.04

%post
    apt-get -y update
    apt-get -y install cowsay lolcat

%environment
    export LC_ALL=C
    export PATH=/usr/games:$PATH

%runscript
    date | cowsay | lolcat
```

*Example source: [https://apptainer.org/docs/user/latest/definition\\_files.html](https://apptainer.org/docs/user/latest/definition_files.html)*





# Now let's build and run a container

```
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer build lolcow.sif lolcow.def
...will take a couple of minutes

# run default container behavior
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer run lolcow.sif

# run your own script from container
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer exec lolcow.sif /bin/sh lolcow.sh
```

# Now let's pull and run a Docker container

```
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer pull cp2k.sif docker://cp2k/cp2k:latest
...will take a couple of minutes

# run at command line:
[monaghaa@c3cpu-a5-u28-1 ~]$ apptainer run cp2k.sif cp2k -o md300.out md300.inp

# schedule a job
[monaghaa@c3cpu-a5-u28-1 ~]$ sbatch run_cp2k.sh
```

# More complex topics (if time allows)

- Bind mounting external directories
- Using containers on GPUs
- Using containers with MPI

# Thank you!

## Survey and feedback

<http://tinyurl.com/curc-survey18>

